

STEERDUPLEX : Steerable Duplex Speech Dialogue Models

Anonymous ACL submission

Abstract

Full-duplex spoken dialogue models are increasingly competent at low-latency turn-taking, interruption handling, and backchanneling, yet a key capability remains under-explored: *steerability*, the ability to reliably shift conversational behavior along target attributes such as tone, persona, speaking speed, voice style, and mid-conversation edits in response to user instructions. We present the first systematic empirical analysis of steerability in current full-duplex models and find that they fall substantially short on both text-based and audio-based steering despite their strong interactive behavior. To close this gap we introduce STEERDUPLEX, a reinforcement-learning post-training method that couples Group Relative Policy Optimization with a *composite reward* mixing (i) audio-aware rubric scores judged jointly over the generated speech and its transcript with (ii) programmatically verifiable signals on objectively measurable attributes. Alongside the method we release STEERBENCH, a hand-graded benchmark of 390 spoken prompts and 1,067 binary audio and text rubrics covering tone, persona, style/accent, and speed/length. Across STEERBENCH, Full-Duplex-Bench v1/v2, VoiceBench, and Audio MultiChallenge, STEERDUPLEX improves average STEERBENCH APR by +22.8 points and Audio MultiChallenge ARS by +7.9 points over the strongest open-source full-duplex baselines while preserving general dialogue quality.¹ Project Page: <https://audioanon.github.io/steerduplex/>.

1 Introduction

Large language models (LLMs) have demonstrated remarkable progress across a broad range of tasks, from natural language understanding and multi-step reasoning to open-ended creative generation (Brown et al., 2020; Bubeck et al., 2023; Guo

et al., 2025). Yet for all their capability, these systems communicate through text, a modality that, while expressive, strips away the richness of human communication. Speech, on the other hand, is the most natural interface through which humans interact, conveying not only content but prosody, emotion, pace, and personality.

This has motivated a growing body of work on spoken dialogue models that operate directly in the speech domain (Nguyen et al., 2023; Veluri et al., 2024; Wang et al., 2024; Zhang et al., 2024; Hu et al., 2025). Operating directly on speech unlocks richer, more natural interaction, letting systems respond to prosody, emotion, and paralinguistic cues in real time. The frontier of this paradigm is the *full-duplex* setting, where a model listens and speaks simultaneously, supporting natural turn-taking, backchanneling, and interruption handling without any explicit segmentation into speaker turns. Our work builds on the publicly available full-duplex backbone of Défossez et al. (2024); the rest of the paper refers to our trained model as STEERDUPLEX.

A key property of any interactive system, however, is that interaction must be steerable. A dialogue system that handles turn-taking and interruptions seamlessly but cannot adjust its tone when asked, or sustain a persona across a conversation, is fundamentally limited in practice. In real interactions, users do not simply want responses; they want responses delivered in a particular way, more formal, slower, in a different voice style. This is a natural part of human conversation, and the ability to reliably shift a model’s behavior along such attributes in response to user instructions is what we term *steerability*. Steerability is already an established concern in text LLMs (Chang et al., 2024; Liang et al., 2024; Keskar et al., 2019), where post-training methods have shown strong gains in persona consistency and style control. Speech, however, makes steerability both more important and

¹Code, model checkpoints, and STEERBENCH will be released upon paper acceptance.




Figure 1: **A capability taxonomy for full-duplex spoken dialogue.** Three pillars: (1) *Text-based steerability* via natural-language instructions (tone, persona, speed/length, mid-conversation edits, graceful refusal); (2) *Audio-based steerability* over acoustic cues (paralinguistic understanding, voice/style editing, audio-context memory, self-coherence); and (3) *General duplex capabilities* (turn-taking, interruption handling, backchanneling, pause handling, mid-utterance corrections). STEERDUPLEX targets all three pillars; STEERBENCH evaluates (1)+(2) and Full-Duplex-Bench v1/v2 (Lin et al., 2025, 2026) evaluates (3).

more challenging: more important because speech carries far richer dimensions of control including prosody, pace, and voice style, and more challenging because these attributes manifest acoustically and cannot be captured or evaluated from text alone.

How steerable, then, are current full-duplex spoken dialogue models? To answer this question, we first introduce STEERBENCH, a benchmark that evaluates steerability along two complementary axes: text-based steerability, covering attributes such as tone, persona, and speaking style, and audio-based steerability, covering paralinguistic understanding, voice and style editing, and self-coherence (Figure 1). Evaluating a range of open and closed full-duplex models on STEERBENCH, we find that existing systems fall substantially short on both axes despite their strong interactive behavior. This motivates the central question of the paper: *can targeted post-training make full-duplex models steerable without degrading their general conversational capability?*

To answer this question constructively, we turn to reinforcement learning. Reinforcement Learning with Verifiable Rewards (RLVR), in which outcome-level rewards are computed directly from

the output rather than from a learned reward model, has emerged as the dominant post-training paradigm in text LLMs (Shao et al., 2024; Guo et al., 2025; Lambert et al., 2024) and has recently generalized to empathetic dialogue (Wang et al., 2025b), multimodal emotion recognition (Zhao et al., 2025), and world-model training (Wu et al., 2025b). RL is therefore the natural candidate; the question is how to design rewards that actually move a duplex model along the steerability axes of Figure 1 without breaking its streaming behavior.

Naively applying RL to steerability in duplex speech runs into two opposing failure modes. Most attributes in Figure 1 (tone, persona, register, voice style) resist clean programmatic verification, so a natural choice is an LLM-as-judge over rubrics (Lambert et al., 2024; Abdulhai et al., 2025; Arora et al., 2026). But *text-only* rubrics miss the prosodic, paralinguistic, and self-coherence axes (Figure 1, pillar 2) that motivate going beyond text dialogue in the first place, and purely rubric-based rewards are also gameable and weak on objectively checkable constraints such as latency, length, language identity, or refusal. STEERDUPLEX therefore optimizes a *composite reward* that combines (a) audio-aware rubrics scored jointly over the gen-

erated speech and its transcript with (b) programmatically verifiable signals on attributes that admit deterministic checks. We find that this composition, rather than any single component, is what produces stable steerability gains on a duplex model.

To our knowledge, no prior work treats steerability as a post-training objective in full-duplex spoken dialogue. ASPIRIN (Hsiao et al., 2026) applies GRPO but only to *when-to-speak* timing, while PERSONAPLEX (Roy et al., 2026) and F-ACTOR (Züfle et al., 2026) achieve narrower control through supervised imitation.

Key contributions. (i) We introduce STEERBENCH, the first benchmark for steerability in full-duplex spoken dialogue, with 390 prompts and 1,067 hand-written rubrics that grade both content and delivery against a per-prompt reference clip across tone, persona, style/accent, and speed/length (§4.6). (ii) We propose STEERDUPLEX, a small-step KL-constrained RL recipe that couples MO-GRPO with a multiplicative composite reward over LLM rubrics (text and reference-anchored audio), verifiable signals (IFEval, safety, latency, FDB timing), and audio-health gates (DNSMOS, silence/clipping/RMS) (§4). (iii) We show via single-family reward probes that no isolated reward family suffices: each probe reaches high scalar reward while collapsing duplex-FDB to zero and leaving 26–49% of rollouts empty, whereas the composite reduces the empty rate to 4% and preserves duplex behavior (§5.3). (iv) STEERDUPLEX improves the strongest open-source duplex baselines by +22.8 points on STEERBENCH APR and +7.9 points on Audio MultiChallenge ARS while remaining competitive on Full-Duplex-Bench v1/v2 and VoiceBench (§5).

2 Related Work

Full-Duplex Spoken Dialogue Models. End-to-end two-channel dialogue modeling was introduced by Nguyen et al. (2023), with subsequent architectures exploring different ways to interleave the two channels in real time: SyncLLM (Veluri et al., 2024), OmniFlatten (Zhang et al., 2024), Wang et al. (2024), and SALM-DUPLEX (Hu et al., 2025). The line culminates in MOSHI (Défossez et al., 2024), which models user and system speech as parallel streams of hierarchical audio tokens with a time-aligned *Inner Monologue* text channel at ~200ms latency. Closest to our setting, Roy et al. (2026) and Züfle et al. (2026) add voice-identity,

role, and interactional control to MOSHI-style models via supervised imitation, but cover only a narrow slice of Figure 1. Our RL recipe operates at the level of sampled rollouts and a paired text channel, so it applies to this entire family of stream-parallel duplex models.

Steerability and Controllable Generation. Controllable text generation, conditioning outputs on attributes such as sentiment, style, persona, and tone, has been studied extensively (Liang et al., 2024; Keskar et al., 2019; Dathathri et al., 2020). Chang et al. (2024) establish steerability as distinct from instruction following, and post-training is known to close the gap: Chen et al. (2024a) find SFT outperforms prompting for personality control, while Abdulhai et al. (2025) use behavioral-consistency metrics as RL signals to reduce persona drift. All such work, however, is text-only and ignores the audio-perceptual and interactional pillars of Figure 1.

RL for Spoken Dialogue. RLVR’s gains in text LLMs (Shao et al., 2024; Guo et al., 2025; Lambert et al., 2024) are now reaching speech: Wu et al. (2025a) apply DPO on AI-annotated preference pairs to MOSHI; Arora et al. (2026) apply multi-reward RLAIIF-DPO to turn-based and duplex systems; and Hsiao et al. (2026) use GRPO for turn-taking timing. None target steerability. We are the first to combine audio-aware rubric rewards, which cover the subjective acoustic and lexical axes that resist programmatic checks, with verifiable rewards that anchor measurable behaviors and limit rubric-judge hacking.

3 Preliminaries

We post-train STEERDUPLEX with on-policy group-relative reinforcement learning. This section reviews the three policy-gradient variants we use (GRPO, GDPO, MO-GRPO) and the composite reward they consume; ablations over algorithm and aggregation choices are reported with the arm grid in Appendix I.

3.1 Group-Relative Policy Optimization

Setup. For each prompt $q \sim \mathcal{D}$, a behavior policy $\pi_{\theta_{\text{old}}}$ samples a group of G rollouts $\{o_i\}_{i=1}^G$, and each rollout receives a scalar reward $R_i = R(o_i; q) \in [0, 1]$. Writing $\mathbf{R} = \{R_i\}_{i=1}^G$ for the group of rewards, $r_{i,t} = \pi_{\theta}(o_{i,t} \mid q, o_{i,<t}) / \pi_{\theta_{\text{old}}}(o_{i,t} \mid q, o_{i,<t})$ for the per-token importance ratio, ε for the clip range, and β for the KL coefficient against a frozen reference π_{ref} , all three policy-gradient variants we consider optimize

the same outer objective:

$$\mathcal{J}(\theta) = \mathbb{E}_{q, \{o_i\}} \left[\frac{1}{G} \sum_{i,t} \min(r_{i,t} \hat{A}_i, \text{clip}(r_{i,t}, 1-\varepsilon, 1+\varepsilon) \hat{A}_i) - \beta D_{\text{KL}}(\pi_\theta \| \pi_{\text{ref}}) \right], \quad (1)$$

where $\hat{A}_i$ is broadcast to every token of rollout i . The variants differ *only* in how $\hat{A}_i$ is computed from $\mathbf{R}$, which is the topic of the next subsection.

GRPO advantage. GRPO (Shao et al., 2024; Guo et al., 2025) computes a single group-normalized advantage from the scalar group reward,

$$\hat{A}_i^{\text{GRPO}} = \frac{R_i - \text{mean}(\mathbf{R})}{\text{std}(\mathbf{R}) + \epsilon}. \quad (2)$$

This works well when R_i is essentially a single quantity, but in our setting $R_i = R(o_i; q)$ aggregates many heterogeneous reward components (§3.3), and the single scalar normalization lets whichever component has the largest within-group variance dominate the gradient.

3.2 GDPO and MO-GRPO

GDPO. Group-Decoupled Normalization Policy Optimization (Li et al., 2026) z -scores each component within the group, then combines with the same human weights w_k as the composite reward:

$$\hat{A}_i^{\text{GDPO}} = \sum_k w_k \cdot \frac{R_i^{(k)} - \text{mean}(\mathbf{R}^{(k)})}{\text{std}(\mathbf{R}^{(k)}) + \epsilon}. \quad (3)$$

A saturated or noisy component now contributes only its standardized residual, so it cannot crowd out signal from the others.

MO-GRPO (default). Multi-objective GRPO (Wang et al., 2025a) uses the same per-component z -scoring as GDPO but adds a $1/\sqrt{K}$ factor for a K -component reward so that the magnitude of the combined advantage is independent of how many components are enabled:

$$\hat{A}_i^{\text{MO-GRPO}} = \frac{1}{\sqrt{K}} \sum_{k=1}^K w_k \cdot \frac{R_i^{(k)} - \text{mean}(\mathbf{R}^{(k)})}{\text{std}(\mathbf{R}^{(k)}) + \epsilon}. \quad (4)$$

Our composite reward mixes binary IFEval verifiers, discrete 1–5 rubric ratings, and continuous DNSMOS scores. Across the sweep in Appendix I, MO-GRPO is the only variant whose update magnitude stays stable as we add or remove reward components, so we adopt it as the default.

3.3 Composite Reward

All three variants above optimize the same composite reward, which is a multiplicative soft gate over reward families:

$$R(o; q) = \prod_k \text{clip}_{[0,1]}(R_k(o; q))^{w_k / \sum_j w_j}. \quad (5)$$

Here R_k ranges over text rubrics, reference-audio rubrics, verifiable instruction checks, safety checks, FDB-style timing checks, promptness, and audio-health gates. The multiplicative form preserves a single scalar reward for the policy-gradient update while making severe failures hard to hide: a rollout that is silent, mistimed, unsafe, or off-instruction cannot be rescued by a high score on any one component. Empirically, this property is what separates valid duplex behavior from superficial reward increases (§5.3).

4 Method

STEERDUPLEX is a post-training recipe for full-duplex speech models, not a new architecture. The method is organized around one constraint: improve steerability across all three pillars of the capability taxonomy in Figure 1 (text-based steering, audio-based steering, and general duplex interaction) without breaking the streaming interface that makes duplex conversation useful. Figure 2 summarizes the pipeline: a Moshi-style backbone receives a consistent steering prompt in its agent-text stream, supervised fine-tuning preserves the duplex interface on high-quality conversational speech, and RL improves steerability through grouped rollouts scored by the composite reward of §3.

4.1 Backbone, Training, and Inference Changes

STEERDUPLEX keeps the model class of Défossez et al. (2024) - three jointly modeled streams (user audio, agent audio, agent text), hierarchical audio tokens, and a time-aligned text channel, and reuses the transformer trunk, audio codec, and depth-former layout without surgery. We prepend a natural-language *steerability system prompt* to the agent-text stream at the start of every dialogue; both audio streams stay silent during the prefix, a delimiter token marks its end, and the same layout is used in SFT, RL, and inference to avoid train/deploy distribution shift. Three training-side changes make the recipe steerability-aware: (i) we mask the cross-entropy loss across the prompt prefix so the optimizer only fits the dialogue region; (ii) we retain




Figure 2: **Overview of STEERDUPLEX.** (A) Starting from a Moshi-style full-duplex backbone, STEERDUPLEX keeps the architecture unchanged and conditions behavior by prepending a natural-language steerability system prompt to the agent-text Inner Monologue stream, with both audio streams silent during the prompt prefix. (B) During post-training, each SteerBench scenario is expanded into $G = 16$ audio-text rollouts, scored by a composite reward, and used to update the trainable policy π_θ with MO-GRPO under an adaptive KL constraint against a frozen reference policy π_{ref} . (C) The reward multiplies audio-aware LLM rubric scores, programmatically verifiable rewards, and audio-health gates, so that critical failures in steerability, safety, timing, or audio quality suppress the total reward and reduce reward-hacking pathways.

the SFT codebook-weighting schedule (first audio codebook $5\times$, non-semantic audio tokens $\times 0.02$, padded text tokens $\times 0.3$) because equal-weight audio tokens otherwise collapse RL into repetition loops (Hsiao et al., 2026); and (iii) we use turn-aware chunking that splits long conversations on alignment boundaries rather than arbitrary frame indices. At inference we sample audio tokens at temperature 0.85 and text tokens at 0.65 with top- k of 200 and 25, capping responses at 12 s and input context at 28 s; these settings are applied uniformly across the benchmarks in §5.

4.2 RL Dataset Construction

RL uses 2,000 curated multi-turn scenarios spanning benchmark-core instruction following, duplex interaction, AudioMC-style memory and self-coherence, text steering, and audio steering. Each scenario contains stitched user audio with timestamps, assistant history, transcripts, category metadata, rubrics, verifiable targets, and, when applicable, a reference audio clip that voices the target style. We sample batches with a stratified sampler so every update includes all major families; Appendix D gives category counts.

4.3 KL-Constrained Small-Step Optimization

RL starts from the SFT checkpoint and uses full-parameter MO-GRPO (Wang et al., 2025a) with

small updates over a longer horizon: 6 prompts per iteration, $G = 16$ rollouts per prompt, 40 iterations, learning rate 2×10^{-7} , PPO clip 0.18, and adaptive KL control around target 0.10. Component rewards are normalized within each rollout group before aggregation, which prevents a high-variance reward family from dominating the update. Hyperparameters and ablation grids are listed in Appendix B and Appendix I.

4.4 Multi-Objective Reward Design

The composite reward has three families: audio-aware rubric rewards that score text and delivery criteria with a frontier multimodal judge; verifiable rewards covering IFEval-style constraints (Zhou et al., 2023), safety refusal patterns, word-count and words-per-minute targets, response latency, and FDB-style timing; and audio-health rewards combining clipping, silence, RMS, minimum-duration gates, and DNSMOS (Reddy et al., 2021). Default aggregation is multiplicative, so a severe failure in any critical family suppresses the total reward instead of being hidden by a strong score elsewhere.

4.5 Reference-Audio Steerability Rewards

A reference-audio clip anchors each audio-steering reward. The judge receives the reference clip, the model-generated clip, and both transcripts, then

scores delivery, prosody, rate, articulation, naturalness, and task success without requiring timbre identity unless the rubric explicitly asks for it—turning subjective style judgments into paired comparisons against a fixed target and reducing the ambiguity of labels such as “cheerful” or “formal.” Exact prompts are in Appendix F.

4.6 Steerability Benchmark

STEERBENCH operationalizes the text-based and audio-based pillars of Figure 1 as an evaluation set. It contains 390 spoken prompts spanning tone, persona, style/accent, and speed/length. Each sample combines a steering request with a concrete task and is graded by 1,067 hand-written binary rubrics: 438 audio rubrics and 629 text rubrics. A sample passes only when all applicable text and audio rubrics pass, preventing content success from masking a failure on the requested steering behavior. Inference is audio-in/audio-out with the same system prompt used during training; text rubrics are judged from transcripts, and audio rubrics are judged against a fixed reference clip generated for the target style.

5 Experiments

5.1 Experimental Setup

Training data. SFT uses 7,024.6 hours of curated synthetic and natural conversational speech, with a 13.7-hour validation split; detailed source statistics are in Appendix P. Word-level timestamps are produced with Qwen3-ForcedAligner-0.6B (Shi et al., 2026); RL uses the 2,000-scenario pack in §4.2.

Evaluation benchmarks. We evaluate steerability with STEERBENCH, duplex interaction with Full-Duplex-Bench v1/v2 (Lin et al., 2025, 2026), general spoken instruction following with VoiceBench (Chen et al., 2024b), and multi-turn speech robustness with Audio MultiChallenge (Gosai et al., 2025). We report FDB v1 in the main paper because it tests whether steering gains preserve interactive behavior; the FDB v2 and VoiceBench tables are deferred to the Extra Results section in the appendix (§A).

Models and baselines. We compare STEERDUPLEX (the model produced by SFT followed by composite-reward RL) against MOSHI (Défossez et al., 2024), PERSONAPLEX (Roy et al., 2026), GPT-4o Realtime (OpenAI, 2025), and Gemini

Live (Google DeepMind, 2025). Baseline details are in Appendix N.

Training details. The SFT stage is full fine-tuning on 8 NVIDIA H100 GPUs with the AdamW optimizer (Loshchilov and Hutter, 2019) at learning rate 1×10^{-5} . The RL stage uses the settings in §4.3: full fine-tuning at 2×10^{-7} , group size $G = 16$, 6 prompts per iteration, 40 iterations, adaptive KL with target 0.10 and $\beta_0 = 0.08$, PPO clip 0.18. The composite reward uses multiplicative aggregation by default. Per-component weights are given in Appendix C; compute and reporting details are in Appendix M.

5.2 Main Results

Tables 1, 2, and 3 in the main paper, together with Tables 4 and 5 in the Extra Results section of the appendix (§A), compare STEERDUPLEX against open- and closed-source baselines. The headline pattern is uniform: we push the open-source frontier substantially closer to proprietary systems on every benchmark targeting steerability, multi-turn robustness, or task-grounded duplex behavior, while preserving the streaming fidelity that motivates full-duplex modeling.

Steerability (Table 1). MOSHI and PERSONAPLEX average only 21.1 and 15.0 APR over the four STEERBENCH axes, versus 53.0 and 57.3 for Gemini Live and GPT-4o Realtime. STEERDUPLEX averages 43.9 APR (+22.8 over the strongest open baseline), with the largest lifts on tone (+33.0) and persona (+37.5) and a substantial gain on style/accent (+28.2). Speed/length is the one axis where MOSHI is already competitive (it tends to emit short responses), and we trade a small amount of length-control APR for stronger gains elsewhere. The audio-side TPR gains confirm that the improvement is not merely tighter text-rubric alignment—it reflects acoustic delivery changes graded against per-prompt reference clips.

Multi-turn robustness (Table 2). On Audio MultiChallenge, STEERDUPLEX lifts overall APR from 6.64 to 10.58 and ARS from 20.56 to 22.19 (+7.9 ARS). On inference memory we reach 20.0 APR, matching Gemini Live’s audio-output number (19.7), and voice editing—on which every prior open-source duplex model scores 0.00—moves off the floor to 1.28. The residual gap to closed systems isolates dimensions audio post-training alone cannot close: voice editing in particular requires

Table 1: **STEERBENCH results.** APR and TPR (% , $\uparrow$) for two proprietary baselines (in grey) and three open-source systems across four steering axes. Best open-source value per column is in **bold**.

Dimension Metric	Tone		Persona		Style & Accent		Speed & Length	
	APR ($\uparrow$)	TPR ($\uparrow$)	APR ($\uparrow$)	TPR ($\uparrow$)	APR ($\uparrow$)	TPR ($\uparrow$)	APR ($\uparrow$)	TPR ($\uparrow$)
<i>Proprietary</i>								
Gemini Live	45.0	68.5	41.0	70.0	73.3	74.7	52.8	76.9
GPT-4o Realtime	49.0	90.6	42.6	93.3	60.4	94.3	77.4	96.2
<i>Open source</i>								
MOSHI (Défossez et al., 2024)	20.0	26.0	5.9	17.6	18.8	21.7	39.6	29.5
PERSONAPLEX (Roy et al., 2026)	10.0	60.6	0.7	39.0	13.4	37.6	35.8	48.7
STEERDUPLEX (ours)	53.0	66.1	43.4	59.2	47.0	46.5	32.1	43.6

Table 2: **Performance comparison on Audio MultiChallenge.** We report overall APR (%) and ARS (%), and per-axis APR across Inference Memory (IM), Instruction Retention (IR), Self Coherence (SC), and Voice Editing (VE). Proprietary baselines are listed in grey for context (numbers reproduced from Gosai et al., 2025, audio-output rows where available). Best open-source value per column is in **bold**.

System	Overall		Per-Axis APR (%)			
	APR (%)	ARS (%)	IM	IR	SC	VE
<i>Proprietary</i>						
Gemini Live (2.5 Flash)	26.11	65.42	19.70	29.17	31.33	26.50
GPT-4o Realtime	20.35	63.03	19.70	20.00	26.51	17.09
<i>Open source</i>						
MOSHI	3.98	14.29	9.09	2.50	3.61	0.00
PERSONAPLEX	6.64	20.56	10.61	6.67	9.64	0.00
STEERDUPLEX (ours)	10.58	22.19	20.00	6.02	15.49	1.28

codec- and decoder-level control beneath the layers our recipe modifies.

Duplex behavior (Table 3; FDB v2 results in Appendix A, Table 4). FDB v1 acts as a guardrail on streaming behavior: STEERDUPLEX obtains the best open-source pause-handling take-over rate (0.372 synthetic, 0.565 Candor) and the best backchannel JSD (0.694), substantially exceeding PERSONAPLEX on every pause-related metric while keeping turn-taking on par with MOSHI. FDB v2 (slow-examiner) is where the composite-reward recipe is most decisive—it requires content correctness *and* delivery simultaneously: we average 4.06 over correction, entity tracking, and safety, surpassing both open baselines (MOSHI 2.78, PERSONAPLEX 1.16) and edging out GPT-4o Realtime on safety (4.29 vs. 4.14).

General spoken capability (Table 5 in Appendix A). STEERDUPLEX improves overall VoiceBench from 38.55 to 46.36 (+7.8), with gains on AlpacaEval, CommonEval, WildVoice, BBH, and AdvBench. IFEval is the one column where we trade accuracy for steering and safety gains—an expected effect of optimizing rubric-graded delivery alongside binary verifiers. The proprietary gap

remains large (−32 points to GPT-4o Realtime), reflecting backbone scale rather than post-training choices.

Why we improve, why open models still trail.

Three choices drive the gains: (i) multiplicative composite rewards keep any single family from carrying the gradient (§5.3); (ii) reference-audio rubric judging turns subjective style adjectives into per-prompt paired comparisons, unlocking the audio TPR gains in Table 1; and (iii) KL-constrained small steps preserve streaming behavior, so FDB v1 stays intact while FDB v2 climbs. The residual 14–41-point gap to GPT-4o Realtime and Gemini Live reflects backbone scale, codec expressiveness, and judge supervision—pointing at joint backbone–codec scaling and richer audio preferences as next steps.

5.3 Reward-Hacking Analysis

The multi-benchmark gains above are *coupled*—only possible because the reward forces every dimension healthy simultaneously. To confirm no single reward family suffices, we train four diagnostic single-family probes and score endpoint rollouts against the full held-out reward battery. Figure 3 and Table 11 report the headline numbers.

The probes expose three failure modes. *Promptness-only* optimization gives the highest scalar reward (0.670) but zero duplex-FDB—precisely the failure mode that would invalidate the FDB v2 gains in Table 4. *Audio-quality-only* raises audio-health and MOS while still scoring zero on duplex-FDB and producing 41/96 empty rollouts, so clean speech can be task-irrelevant. *Text-rubric-only* and *audio-rubric-only* probes produce 47/96 and 41/96 empty rollouts, showing judge rewards must be guarded by transcript and audio validity. The joint composite keeps the empty rate at 4/96 and lifts duplex-FDB to 0.371 at a scalar reward below the promptness probe but representing valid

Table 3: **Performance comparison on Full-Duplex-Bench v1 (Lin et al., 2025)**. We report take-over rate (TOR), backchannel frequency (Freq) and JSD against human distributions, and first-response latency across the four FDB-v1 dimensions: Pause Handling, Backchanneling, Smooth Turn Taking, and User Interruption. Arrows indicate whether lower ($\downarrow$) or higher ($\uparrow$) is better. Proprietary baselines are in grey for context; best open-source value per column is in **bold**.

Dimension	Pause Handling		Backchannel			Smooth Turn Taking		User Interruption		
Data Metric	Synthetic	Candor	ICC			Candor		Synthetic		
Model	TOR ($\downarrow$)	TOR ($\downarrow$)	TOR ($\downarrow$)	Freq ($\uparrow$)	JSD ($\downarrow$)	TOR ($\uparrow$)	Latency ($\downarrow$)	TOR ($\uparrow$)	GPT-4o ($\uparrow$)	Latency ($\downarrow$)
<i>Proprietary</i>										
Gemini Live	0.000	0.023	0.000	0.026	0.748	0.008	0.000	0.185	0.135	1.215
GPT-4o Realtime	0.007	0.040	0.000	0.008	0.974	0.028	0.000	0.098	0.125	2.615
<i>Open source</i>										
MOSHI	1.000	1.000	0.455	0.164	0.708	1.000	0.000	0.895	3.687	0.973
PERSONAPLEX	0.664	0.676	0.527	0.073	0.765	0.975	0.157	0.945	4.566	0.257
STEERDUPLEX (ours)	0.372	0.565	0.564	0.125	0.694	0.739	0.211	0.753	3.732	0.970




Figure 3: **Reward-hacking diagnostics for isolated reward families**. At the diagnostic endpoint, four single-family probes (text-rubric, audio-rubric, audio-quality, promptness) and the joint composite are scored on a held-out diagnostic battery. (a) Scalar reward rises across probes but empty rate stays high and duplex-FDB collapses to zero on every isolated probe; only the joint composite reaches a non-zero duplex score with a near-zero empty rate. (b) Each probe maxes out only the components related to its own family while the *duplex* column (red outline) is uniformly near zero, confirming that duplex behavior is the shared failure mode of isolated rewards. (c) Plotting scalar reward against empty rate places every isolated probe in the upper reward-hack region, while the joint composite (green star) lands inside the target zone.

behavior across every component.

6 Conclusion and Future Work

The transferable takeaway is that a scalar reward curve is insufficient evidence of improved duplex behavior unless monitored alongside transcript validity, timing, content, safety, and audio-health diagnostics. Multiplicative aggregation is the mechanism that carries this property into the optimizer—a severe failure in any required family suppresses the total reward rather than being averaged away—and is also what drives the simultaneous gains in Tables 1, 2, and 3 (with Tables 4 and 5 in Appendix A). Any duplex post-training pipeline targeting style or steerability should pair its judge signals with verifiable and audio-health gates, or risk scoring well on reward dashboards while shipping unusable conversations. Component traces are in Appendix J.

We presented STEERDUPLEX, a post-training recipe that improves steerability in full-duplex speech models without changing the base architecture, combining KL-bounded small-step RL with a composite reward over reference-audio rubrics, text rubrics, verifiable checks, and audio-health gates. Across STEERBENCH, Audio MultiChallenge, FDB v1/v2, and VoiceBench, it lifts the strongest open-source baselines by +22.8 APR on steerability and +7.9 ARS on multi-turn robustness while preserving duplex interaction, and our reward-hacking probes show that no isolated reward family suffices, making steerability a concrete post-training target for future open duplex models.

Limitations

STEERBENCH focuses on controlled English steering requests with fixed reference clips, and does not cover all languages, dialects, recording conditions, or long-horizon personalization. The audio-rubric judge depends on frontier multimodal evaluation, so absolute scores should be read alongside transcript validity, audio-health diagnostics, and downstream benchmarks. The model also inherits the latency, codec, and single-voice constraints of the Moshi-style backbone, bounding how far post-training alone can move the interaction frontier.

References

Marwa Abdulhai, Ryan Cheng, Donovan Clay, Tim Althoff, Sergey Levine, and Natasha Jaques. 2025. Consistently simulating human personas with multi-turn reinforcement learning. In *Advances in Neural Information Processing Systems*.

Siddhant Arora, Jinchuan Tian, Jiatong Shi, Hayato Futami, Yosuke Kashiwagi, Emiru Tsunoo, and Shinji Watanabe. 2026. Optimizing conversational quality in spoken dialogue systems with reinforcement learning from AI feedback. *arXiv preprint arXiv:2601.19063*.

Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, and 1 others. 2020. Language models are few-shot learners. In *Advances in Neural Information Processing Systems*, volume 33, pages 1877–1901.

Sébastien Bubeck, Varun Chandrasekaran, Ronen Eldan, Johannes Gehrike, Eric Horvitz, Ece Kamar, Peter Lee, Yin Tat Lee, Yuanzhi Li, Scott Lundberg, and 1 others. 2023. Sparks of artificial general intelligence: Early experiments with GPT-4. *arXiv preprint arXiv:2303.12712*.

Trenton Chang, Tobias Schnabel, Adith Swaminathan, and Jenna Wiens. 2024. Measuring steerability in large language models. In *NeurIPS Workshop on Safe Generative AI*.

Yanquan Chen, Zhen Wu, Junjie Guo, Shujian Huang, and Xinyu Dai. 2024a. Extroversion or introversion? controlling the personality of your large language models. *arXiv preprint arXiv:2406.04583*.

Yiming Chen, Xianghu Yue, Chen Zhang, Xiaoxue Gao, Robby T. Tan, and Haizhou Li. 2024b. VoiceBench: Benchmarking LLM-based voice assistants. *arXiv preprint arXiv:2410.17196*.

Christopher Cieri, David Miller, and Kevin Walker. 2004. The Fisher corpus: A resource for the next generations of speech-to-text. *Proc. LREC*.

Sumanth Dathathri, Andrea Madotto, Janice Lan, Jane Hung, Eric Frank, Piero Molino, Jason Yosinski, and Rosanne Liu. 2020. Plug and play language models: A simple approach to controlled text generation. In *International Conference on Learning Representations*.

Alexandre Défossez, Laurent Mazaré, Manu Orsini, Amélie Royer, Patrick Pérez, Hervé Jégou, Edouard Grave, and Neil Zeghidour. 2024. Moshi: A speech-text foundation model for real-time dialogue. *arXiv preprint arXiv:2410.00037*.

Google DeepMind. 2025. Gemini Live API: Real-time multimodal dialogue. Google AI for Developers documentation. Accessed via the gemini-3.1-flash-live-preview API endpoint.

Advait Gosai, Tyler Vuong, Utkarsh Tyagi, Steven Li, Wenjia You, Miheer Bavare, Arda Uçar, Zhongwang Fang, Brian Jang, Bing Liu, and Yunzhong He. 2025. Audio MultiChallenge: A multi-turn evaluation of spoken dialogue systems on natural human interaction. *arXiv preprint arXiv:2512.14865*.

Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shitong Ma, Peiyi Wang, Xiao Bi, and 1 others. 2025. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint arXiv:2501.12948*.

Chi-Yuan Hsiao, Ke-Han Lu, Yu-Kuan Fu, Guan-Ting Lin, Hsiao-Tsung Hung, and Hung-yi Lee. 2026. ASPIRin: Action space projection for interactivity-optimized reinforcement learning in full-duplex speech language models. *arXiv preprint arXiv:2604.10065*.

Ke Hu, Ehsan Hosseini-Asl, Chen Chen, Edresson Casanova, Subhankar Ghosh, Piotr Żelasko, Zhehuai Chen, Jason Li, Jagadeesh Balam, and Boris Ginsburg. 2025. SALM-Duplex: Efficient and direct duplex modeling for speech-to-speech language model. In *Proc. Interspeech*.

Nitish Shirish Keskar, Bryan McCann, Lav R. Varshney, Caiming Xiong, and Richard Socher. 2019. CTRL: A conditional transformer language model for controllable generation. *arXiv preprint arXiv:1909.05858*.

Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, Faeze Brahman, Lester James V. Miranda, Alisa Liu, Nouha Dziri, Shane Lyu, and 1 others. 2024. Tulu 3: Pushing frontiers in open language model post-training. *arXiv preprint arXiv:2411.15124*.

Hao Li, Chen Zhang, and 1 others. 2026. GDPO: Group reward-decoupled normalization for policy optimization in multi-reward RLHF. *arXiv preprint arXiv:2601.05242*.

Xun Liang, Hanyu Wang, Yezhaohui Wang, Shichao Song, Jiawei Yang, Simin Niu, Jie Hu, Dan Liu,

679	Shunyu Yao, Feiyu Xiong, and Zhiyu Li. 2024. Controllable text generation for large language models: A survey. <i>arXiv preprint arXiv:2408.12599</i> .	
680		
681		
682	Guan-Ting Lin, Shih-Yun Shan Kuan, Jiatong Shi, Kai-Wei Chang, Siddhant Arora, Shinji Watanabe, and Hung-yi Lee. 2026. Full-Duplex-Bench-v2: A multi-turn evaluation framework for duplex dialogue systems with an automated examiner. <i>arXiv preprint arXiv:2510.07838</i> .	
683		
684		
685		
686		
687		
688	Guan-Ting Lin, Jiachen Lian, Tingle Li, Qirui Wang, Gopala Anumanchipalli, Alexander H. Liu, and Hung-yi Lee. 2025. Full-Duplex-Bench: A benchmark to evaluate full-duplex spoken dialogue models on turn-taking capabilities. In <i>Proc. ASRU</i> .	
689		
690		
691		
692		
693	Ilya Loshchilov and Frank Hutter. 2019. Decoupled weight decay regularization. In <i>International Conference on Learning Representations</i> .	
694		
695		
696	Tu Anh Nguyen, Eugene Kharitonov, Jade Copet, Yossi Adi, Wei-Ning Hsu, Ali Elkahky, Paden Tomasello, Robin Algayres, Benoît Sagot, Abdelrahman Mohamed, and Emmanuel Dupoux. 2023. Generative spoken dialogue language modeling. <i>Transactions of the Association for Computational Linguistics</i> , 11:250–266.	
697		
698		
699		
700		
701		
702		
703	OpenAI. 2025. Introducing the Realtime API (GPT-4o realtime). OpenAI Platform documentation. Accessed via the <code>gpt-realtime</code> API endpoint.	
704		
705		
706	Chandan K.A. Reddy, Vishak Gopal, and Ross Cutler. 2021. DNSMOS: A non-intrusive perceptual objective speech quality metric to evaluate noise suppressors. In <i>IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)</i> .	
707		
708		
709		
710		
711	Rajarshi Roy, Jonathan Raiman, Sang-gil Lee, Teodor-Dumitru Ene, Robert Kirby, Sungwon Kim, Jaehyeon Kim, and Bryan Catanzaro. 2026. PersonaPlex: Voice and role control for full duplex conversational speech models. <i>arXiv preprint arXiv:2602.06053</i> .	
712		
713		
714		
715		
716	Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y.K. Li, Y. Wu, and Daya Guo. 2024. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. <i>arXiv preprint arXiv:2402.03300</i> .	
717		
718		
719		
720		
721		
722	Xian Shi, Xiong Wang, Zhifang Guo, Yongqi Wang, Pei Zhang, Xinyu Zhang, Zishan Guo, Hongkun Hao, Yu Xi, Baosong Yang, Jin Xu, Jingren Zhou, and Junyang Lin. 2026. Qwen3-ASR technical report. <i>arXiv preprint arXiv:2601.21337</i> .	
723		
724		
725		
726		
727	Bandhav Veluri, Benjamin N. Peloquin, Bokai Yu, Hongyu Gong, and Shyamnath Gollakota. 2024. Beyond turn-based interfaces: Synchronous LLMs as full-duplex dialogue agents. In <i>Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing</i> , pages 21390–21402.	
728		
729		
730		
731		
732		
	Hanyang Wang, Lei Sun, Zichen Wei, and 1 others. 2025a. MO-GRPO: Stabilizing multi-objective reinforcement learning with reward normalization. <i>arXiv preprint arXiv:2509.22047</i> .	733
		734
		735
		736
	Peisong Wang, Ruotian Ma, Bang Zhang, Xingyu Chen, Zhiwei He, Kang Luo, Qingsong Lv, Qingxuan Jiang, Zheng Xie, Shanyi Wang, Yuan Li, Fanghua Ye, Jian Li, Yifan Yang, Zhaopeng Tu, and Xiaolong Li. 2025b. RLVER: Reinforcement learning with verifiable emotion rewards for empathetic agents. <i>arXiv preprint arXiv:2507.03112</i> .	737
		738
		739
		740
		741
		742
		743
	Peng Wang, Songshuo Lu, Yaohua Tang, Sijie Yan, Wei Xia, and Yuanjun Xiong. 2024. A full-duplex speech dialogue scheme based on large language models. In <i>Advances in Neural Information Processing Systems</i> .	744
		745
		746
		747
	Anne Wu, Laurent Mazaré, Neil Zeghidour, and Alexandre Défossez. 2025a. Aligning spoken dialogue models from user interactions. In <i>International Conference on Machine Learning</i> .	748
		749
		750
		751
	Jialong Wu, Shaofeng Yin, Ningya Feng, and Mingsheng Long. 2025b. RLVR-World: Training world models with reinforcement learning. In <i>Advances in Neural Information Processing Systems</i> .	752
		753
		754
		755
	Qinglin Zhang, Luyao Cheng, Chong Deng, Qian Chen, Wen Wang, Siqi Zheng, Jiaqing Liu, Hai Yu, Chaofeng Tan, Zhihao Du, and 1 others. 2024. OmniFlatten: An end-to-end GPT model for seamless voice conversation. <i>arXiv preprint arXiv:2410.17799</i> .	756
		757
		758
		759
		760
		761
	Jiaxing Zhao, Xihan Wei, and Liefeng Bo. 2025. R1-Omni: Explainable omni-multimodal emotion recognition with reinforcement learning. <i>arXiv preprint arXiv:2503.05379</i> .	762
		763
		764
		765
	Jeffrey Zhou, Tianjian Lu, Swaroop Mishra, Siddhartha Brahma, Sujoy Basu, Yi Luan, Denny Zhou, and Le Hou. 2023. Instruction-following evaluation for large language models. In <i>Workshop on Instruction Tuning and Instruction Following at NeurIPS</i> .	766
		767
		768
		769
		770
	Maike Züfle, Ondrej Klejch, Nicholas Sanders, Jan Niehues, Alexandra Birch, and Tsz Kin Lam. 2026. F-Actor: Controllable conversational behaviour in full-duplex models. <i>arXiv preprint arXiv:2601.11329</i> .	771
		772
		773
		774
		775

Appendix overview

The appendix groups material by the main-paper section it supports:

METHOD (§4). Hyperparameters (§B), reward weights (§C), RL-pack category counts (§D), training and inference prompts (§E), reference-audio rubric judge prompt (§F), per-criterion text rubric judge prompt (§G), and benchmark judge prompt (§H).

EXPERIMENTS (§5). Extra results: FDB v2 and VoiceBench tables (§A), experiment-matrix arm grids (§I), reward-hacking diagnostic components (§J), evaluation benchmark coverage (§K), AudioMC task and speaker coverage (§L), and compute budget (§M).

BASELINES AND QUALITATIVE EXAMPLES. Baseline system descriptions (§N) and qualitative win cases (§O).

DATA AND ETHICS. SFT data statistics (§P), contributor/voice/PII protocol (§Q), demographic coverage (§R), artifact licenses (§S), ethical risks (§T), and AI writing assistance disclosure (§U).

A Extra Results

This section collects supplementary result tables referenced from the main paper. Table 4 reports Full-Duplex-Bench v2 (slow-examiner) scores discussed in §5.2, and Table 5 reports the full VoiceBench breakdown supporting the general-capability paragraph in §5.2.

Table 4: **Performance comparison on Full-Duplex-Bench v2 (Lin et al., 2026)** (slow-examiner setting). Scores are LLM-judged 1–5 ratings on three task families: Correction handling, Entity tracking, and Safety. Proprietary baselines are in grey for context; best open-source value per column is in **bold**.

System	Correction	Entity	Safety
<i>Proprietary</i>			
Gemini Live	3.81	3.53	3.83
GPT-4o Realtime	4.24	4.17	4.14
<i>Open source</i>			
MOSHI (Défossez et al., 2024)	2.90	2.65	2.79
PERSONAPLEX (Roy et al., 2026)	1.13	1.00	1.35
STEERDUPLEX (ours)	4.02	3.87	4.29

B Training Hyperparameters

Table 6 lists every hyperparameter used by the default STEERDUPLEX run, including the SFT opti-

mizer, the MO-GRPO setup, the adaptive-KL controller, codebook weighting, mixed-precision settings, and the inference-side sampling parameters.

C Reward Weights

Table 7 shows the per-component weights w_k that the joint-full configuration uses inside the multiplicative composite reward of Eq. (5). Ablation arms in §5.3 zero out the corresponding rows.

D SteerBench RL Pack: Category Counts

Table 8 reports the per-category sample counts for the 2,000-scenario RL pack used in all our experiments. The pack covers five families (benchmark core, AudioMC-style multi-turn, duplex/FDB, text-based steering, audio-based steering) so that stratified sampling exposes every family in every RL iteration.

E Training, Inference, and Context Prompts

This section lists the three prompt templates used by STEERDUPLEX: the constant training system prompt, the shared inference system prompt, and the multi-turn context wrapper that surfaces dialogue history.

The training-time system prompt below is prepended to the agent-text channel of every RL rollout; per-scenario steering instructions are inserted on the user side via the dialogue history.

You are a helpful full-duplex voice assistant. Your voice and identity are fixed. Stay kind, safe, and constructive. Follow the user’s spoken instructions about tone, persona flavor, speaking style, speed, and length when those instructions are allowed.

For inference and benchmark evaluation, the shared default assistant prompt is:

You are a helpful voice assistant. Your voice and identity are fixed. Listen carefully, follow the user’s instructions, and respond naturally.

For multi-turn RL scenarios rendered as a stitched user-audio stream, prior assistant turns are added to the system prompt with the following context wrapper:

Conversation so far (you produced the assistant lines):
 User: <previous_user_text>
 You: <previous_assistant_text>
 ...
 User (now): <latest_user_text>

Now respond to the user’s latest turn, taking the prior context into account.

Table 5: **General-capability evaluation on VoiceBench (Chen et al., 2024b)** across instruction following, reasoning, safety, and spoken QA. GPT-score based evaluations are reported on a 1–5 scale; SD-QA reports PEDANT / GPT-Judge scores; higher is better for all columns. Proprietary baselines are in grey for context; best open-source value per column is in **bold**.

Model	AlpacaEval	CommonEval	WildVoice	SDQA	IFEval	BBH	AdvBench	Overall
<i>Proprietary</i>								
Gemini Live	3.69	3.52	3.26	53.5 / 40.6	17.0 / 25.9	62.8	94.1	63.01
GPT-4o Realtime	4.65	4.33	4.38	88.9 / 75.9	24.3 / 37.3	62.9	98.8	78.42
<i>Open source</i>								
MOSHI (Défossez et al., 2024)	2.15	1.94	1.49	25.4 / 35.3	9.9 / 19.4	49.9	68.2	38.55
PERSONAPLEX (Roy et al., 2026)	2.45	2.29	1.56	20.6 / 27.8	9.0 / 17.2	48.6	5.3	30.51
STEERDUPLEX (ours)	2.94	2.56	2.26	31.6 / 28.8	6.1 / 12.2	52.8	75.6	46.36

Setting	Value
Base model	MOSHI (Défossez et al., 2024)
SFT optimizer	AdamW, lr 1×10^{-5}
SFT data	7,024.6 h train / 13.7 h validation
RL data	2,000 scenarios, 21 categories
RL algorithm (default)	MO-GRPO (Wang et al., 2025a)
RL mode	Full fine-tuning
Iterations	40
Prompts per iteration	6 (stratified)
Rollouts per prompt G	16
Optimizer steps per iter (max)	32
Learning rate	2×10^{-7}
Weight decay	0.05
Gradient clip norm	0.5
PPO clip ϵ	0.18
KL target	0.10
KL controller β_0	0.08
KL controller bounds	[0.02, 0.8]
KL early-stop multiplier	2.0
First codebook multiplier	$5 \times$
Mixed precision	bfloat16
Gradient checkpointing	on
Sampling temperature (audio)	0.85
Sampling temperature (text)	0.65
Top- k (audio / text)	200 / 25
Max response length	12 s

Table 6: Hyperparameters for the default STEERDUPLEX run. Algorithm and KL target are swept in Appendix I.

F Reference-Audio Rubric Judge Prompt

This section reproduces the prompt template used by the reference-audio rubric judge of §4.5. The judge receives two audio clips (Audio 1 the target style reference, Audio 2 the model output) plus both transcripts, and returns per-criterion JSON scores. The template below is lightly simplified for presentation; the exact string is in the released code.

You are judging audio steerability for a full-duplex voice assistant.
 Audio 1 is the reference target.
 Audio 2 is the model-generated answer.
 Compare delivery / style / prosody / speaking rate / articulation / affect against the reference. Do NOT require identical speaker timbre or voice

Component	Weight
Text rubric judge / total	0.22
Audio rubric judge / total (ref-anchored)	0.14
Audio quality / health gate	0.12
Audio quality / MOS	0.08
Duplex-FDB / total (Lin et al., 2025)	0.12
STEERBENCH delivery / total	0.12
Respond-promptly / total	0.08
IFEval / total (Zhou et al., 2023)	0.06
Safety / total	0.06

Table 7: Default composite-reward weights w_i for the joint-full arm. Weights are normalized when used in the multiplicative aggregator. Ablation arms zero out the corresponding rows.

identity unless the rubric explicitly asks for it. Use the transcript only to verify content; judge audio qualities from the audio.

Scenario category: <category>
 User request/context: <user_text>
 Reference transcript: <ref_text>
 Generated transcript: <gen_text>

Audio rubrics:

<numbered list of rubric items>

Return ONLY valid JSON with numeric scores in [0,1]:

```
{
  "reason": "<one sentence>",
  "score": 0.0,
  "style_match": 0.0,
  "prosody_match": 0.0,
  "speech_rate_length": 0.0,
  "articulation": 0.0,
  "naturalness": 0.0,
  "task_success": 0.0,
  "criteria": [
    {
      "title": "...", "category": "...",
      "rating": "No Issues|Minor Issues|Major Issues",
      "score": 0.0, "reason": "..."
    }
  ]
}
```


Category	Count
<i>Benchmark core</i>	
Daily conversation	150
Safety	140
Correction handling	140
IFEval	120
Entity tracking	120
Factual	90
<i>AudioMC-style multi-turn</i>	
Voice editing	150
Instruction retention	130
Self coherence	130
Inference memory	130
<i>Duplex / FDB</i>	
Pause handling	80
Interruption	50
Backchannel	40
Turn taking	30
<i>Text-based steering</i>	
Dynamic mid-conv steering	100
Style / accent	70
Persona	70
Tone	60
Speed / length	40
<i>Audio-based steering</i>	
Paralinguistic understanding	80
Speech understanding from acoustics	80
Total	2,000

Table 8: Per-category counts in the RL pack used for all experiments. Stratified sampling guarantees every *family* is represented in every batch.

G Per-Criterion Text Rubric Judge Prompt

The text rubric judge of §4.4 is queried once per rubric criterion attached to a scenario, returning a three-level rating with a one-sentence explanation. The template below is what we send for each criterion.

Rate how well this single rubric criterion is satisfied. Evaluate ONLY this criterion.

Rating scale:

- "No Issues": Meets the criterion excellently.
- "Minor Issues": Mostly meets it; small gaps or imperfect phrasing.
- "Major Issues": Significant deficiency or failure on this criterion.

Rules:

- "Explicit" criteria must be directly addressed. "Implicit" may be inferred from context.
- "Objective" = factual pass/fail.
- "Subjective" = quality judgment.
- The "weight" field is context only - it

does NOT change your rating.

- Minor phrasing/formatting differences are OK if substance is correct.
- For audio assistants: judge the TRANSCRIPT - content correctness > delivery polish.

Criterion:

- Title: <title>
- Category: <category>
- <Explicit|Implicit> | <Objective|Subjective> | Weight: <weight>
- Description: <description>

User request (last turn): <user_text>

Assistant transcript: <response>

Return ONLY valid JSON.

"reasoning" BEFORE "rating".

```
{
  "reasoning": "<one sentence>",
  "rating": "No Issues" |
             "Minor Issues" |
             "Major Issues"
}
```

H Benchmark Judge System Prompt

For benchmark-aligned scenarios (daily, correction, safety, entity-tracking, factual, AudioMC), the benchmark judge is conditioned with the following system prompt and a per-task rubric (one of: daily, correction, entity-tracking, safety, alpaca, factual, ifeval, default):

You judge spoken-assistant transcripts. Given the user's request, any scenario-specific rubric, and the assistant's transcript, FIRST write a ONE-SENTENCE explanation of how the assistant performed (referencing concrete details from the transcript), THEN give a 1..5 integer score (5 = best). If a list of criteria is provided, evaluate each one independently with a one-sentence reason and a boolean flag.

Return ONLY valid JSON with this exact key ORDER (reasoning first, scores last) and no other text:

```
{
  "reason": "<ONE sentence, <= 220 chars>",
  "criteria_reasons": ["...", ...],
  "criteria_met":      [<bool>, ...],
  "task_score": <int 1..5>
}
```

Be calibrated. A typical OK-but-imperfect response is a 3. A 5 requires the response to actually accomplish the task as a competent human would. Always reason FIRST, score AFTER - never invert the order.

I Experiment-Matrix Arm Grids

Table 9 enumerates the algorithm-and-KL sweep (six arms varying MO-GRPO / GDPO and the KL

target across 0.05–0.20) and Table 10 enumerates the reward-composition sweep (eight arms toggling reward families, including the four single-reward arms that drive the reward-hacking analysis of §5.3).

Arm	Algorithm	KL target
A1 (default)	MO-GRPO	0.10
A2	MO-GRPO	0.05
A3	MO-GRPO	0.15
A4	MO-GRPO	0.20
A5	GDPO	0.10
A6 (sum)	MO-GRPO	0.10

Table 9: Algorithm-and-KL sweep. A6 differs from A1 in using weighted-sum aggregation instead of multiplicative.

Arm	Enabled rewards
B1 joint-full	all of: text-rubric, audio-rubric, audio-health, MOS, duplex-FDB, STEERBENCH, IFEval, safety, promptness
B2 no-audio-rubric	B1 minus audio-rubric
B3 no-audio-quality	B1 minus health gate and MOS
B4 no-text-rubric	B1 minus text-rubric
B5 rubric-only	text-rubric only
B6 audio-rubric-only	audio-rubric only
B7 audio-quality-only	health gate + MOS only
B8 promptness-only	promptness only

Table 10: Reward-composition ablations. B5–B8 are the single-reward arms used in the reward-hacking analysis of §5.3.

J Reward-Hacking Diagnostic Components

Table 12 expands Table 11 from §5.3 with the full per-component breakdown for each single-family probe (text rubric, audio rubric, audio quality, promptness). Each row reports the scalar composite, the empty-rollout count, and the mean held-out score across all nine reward components, so that the column-wise failure pattern in Figure 3(b) can be read off directly.

K Evaluation Benchmark Coverage

This section gives a one-paragraph description of each evaluation benchmark used in §5 (STEERBENCH, Audio MultiChallenge, Full-Duplex-Bench v1/v2, and VoiceBench) so the role of each benchmark in the headline tables is unambiguous.

Table 11: **Single-family reward probes vs. the joint composite.** Endpoint diagnostics on a held-out set of 96 sampled rollouts. Higher scalar reward does not imply usable behavior: every isolated probe leaves the empty-rollout rate $\geq 26\%$ and the duplex-FDB score at zero.

Probe	Scalar $\uparrow$	Empty/96 $\downarrow$	Duplex $\uparrow$	MOS $\uparrow$
Text-rubric only	0.329	47	0.055	0.512
Audio-rubric only	0.336	41	0.000	0.571
Audio-quality only	0.522	41	0.000	0.561
Promptness only	0.670	25	0.000	0.626
STEERDUPLEX (joint)	0.554	4	0.371	0.631

STEERBENCH. STEERBENCH is our steerability benchmark for full-duplex spoken dialogue. It covers tone, persona, style/accent, and speed/length requests with separate text and audio rubrics, and it evaluates whether a spoken assistant satisfies both the requested task and the requested delivery style.

Audio MultiChallenge (Gosai et al., 2025). Audio MultiChallenge evaluates multi-turn spoken robustness along inference memory, instruction retention, self-coherence, and voice editing. We use it as the main guardrail for whether steerability improvements preserve longer-context speech interaction.

Full-Duplex-Bench v1 (Lin et al., 2025). Full-Duplex-Bench v1 evaluates low-level full-duplex interaction behavior, including pause handling, backchanneling, smooth turn taking, and user interruption. It directly tests whether a model can listen and speak under timing-sensitive streaming conditions.

Full-Duplex-Bench v2 (Lin et al., 2026). Full-Duplex-Bench v2 extends duplex evaluation to multi-turn task settings. We report the slow-examiner setting over correction handling, entity tracking, and safety, which probes whether the model can preserve task state while interacting naturally.

VoiceBench (Chen et al., 2024b). VoiceBench evaluates general spoken instruction-following and knowledge capability across AlpacaEval, CommonEval, WildVoice, SD-QA, IFEval, BBH, and AdvBench. We include it to verify that speech-steering post-training does not merely improve style or timing while sacrificing general problem-solving and safety behavior.

Table 12: **Component diagnostics for single-family reward probes.** Values are means over sampled rollouts at the diagnostic endpoint. A dash indicates that the component was not applicable to the sampled category mix.

Probe	Scalar	Empty/96	Text	Audio	Prompt	Health	MOS	Duplex	IFEval	Safety
Text rubric only	0.329	47	0.324	0.211	0.186	0.521	0.512	0.055	0.391	–
Audio rubric only	0.336	41	0.195	0.259	0.293	0.604	0.571	0.000	0.188	0.125
Audio quality only	0.522	41	0.186	0.278	0.320	0.583	0.561	0.000	0.188	0.125
Promptness only	0.670	25	0.282	0.283	0.532	0.750	0.626	0.000	0.438	–

L AudioMC Task and Speaker Coverage

The Audio MultiChallenge task sheets contain 19,568 multi-turn tasks and 1,000.1 hours of user audio across 2,998 user speakers. The task axis is stored in the `axis` field, while speaker gender and native accent are sheet columns F and G (`user_speaker_gender` and `user_speaker_native_accent`). Table 13 summarizes the axis-level distribution; the metadata shows a balanced gender mix and broad native-accent coverage rather than a single-speaker or single-accent population.

M Compute Budget and Result Reporting

All open-weight models in this work use a Mosh-style 7B-parameter backbone; GPT-4o Realtime and Gemini Live are closed systems whose parameter counts are not disclosed by their providers. The default SFT run uses 8 NVIDIA H100 80GB GPUs, bfloat16 precision, gradient checkpointing, and 31,020 planned optimizer steps; at measured throughput this corresponds to roughly 1.0–1.1k H100 GPU-hours. The default RL run uses full fine-tuning on 8 H100 80GB GPUs with 40 iterations, 6 prompts per iteration, and 16 rollouts per prompt; the reward-ablation suite uses the same hardware profile with shorter diagnostic runs. Unless otherwise stated, reported benchmark numbers are means over 3 independent runs with fixed data splits and different decoding/evaluation seeds; we do not report the maximum over runs.

N Baseline Systems

This section describes the four baselines compared with STEERDUPLEX in §5: MOSHI (open-source backbone), PERSONAPLEX (closest control-focused open baseline), and the two proprietary API systems GPT-4o Realtime and Gemini Live, together with the exact checkpoints and access protocol we use.

MOSHI (Défossez et al., 2024). MOSHI is a 7B-parameter open-source full-duplex spoken dialogue model from Kyutai and the seminal open-weight model in this space. Starting from a text language-model backbone, it generates speech as tokens from the residual quantizer of a neural audio codec and models its own speech and the user’s as parallel streams, removing explicit speaker turns. A small Depth Transformer models inter-codebook dependencies at each step and a 7B Temporal Transformer models temporal dependencies; the theoretical latency is 160 ms and the practical latency is ~ 200 ms. The model is trained on 170 h of natural and scripted multi-speaker conversations and 20,000 h of synthetic dialogue. MOSHI is the base model for STEERDUPLEX and is included as a primary baseline; we use the publicly released *moshiko-pytorch-bf16* checkpoint.

PERSONAPLEX (Roy et al., 2026). PERSONAPLEX is a 7B-parameter full-duplex speech-to-speech dialogue model designed for hybrid system prompts that combine role conditioning via text with voice cloning via speech samples, trained on a large-scale synthetic dataset of paired prompts and user–agent conversations. The model is built on MOSHI weights and uses Helium as the underlying language-model backbone, with a Mimi speech encoder and decoder operating at 24 kHz. Before the conversation, PERSONAPLEX is conditioned on a voice prompt that establishes target vocal characteristics and a text prompt that specifies persona attributes such as role, background, and scenario context. We use the publicly released *personaplex-7b-v1* checkpoint and access it through the official inference server.

GPT-4o Realtime (OpenAI, 2025). GPT-4o Realtime is a closed, API-only speech-to-speech dialogue model from OpenAI, part of the GPT-4o family. Unlike cascaded systems, the model consumes raw audio and produces raw audio over a single persistent WebSocket session, with no transcription

Axis	Tasks	Hours	Turns	Speakers	Gender split	Most common native accents
Inference Memory	5,184	285.6	6.66	1,984	F 52.1 / M 46.4 / O 1.5	General American 30.3%; African (non-SA) 14.0%; British 13.3%; Canadian 9.9%
Instruction Retention	5,352	231.4	6.88	2,019	F 51.6 / M 46.7 / O 1.7	General American 29.8%; African (non-SA) 14.0%; British 12.6%; Canadian 10.1%
Self Coherence	3,880	199.8	7.06	1,695	F 50.5 / M 47.9 / O 1.6	General American 27.0%; African (non-SA) 16.1%; British 15.7%; Canadian 8.5%
Voice Editing	5,152	283.4	7.08	1,930	F 50.8 / M 47.7 / O 1.5	General American 27.3%; British 16.9%; African (non-SA) 15.9%; Canadian 8.8%
Total	19,568	1,000.1	6.91	2,998	F 51.3 / M 47.1 / O 1.6	General American 28.7%; African (non-SA) 14.9%; British 14.5%; Canadian 9.4%

Table 13: AudioMC task and speaker coverage parsed from the two task sheets in `audiomc_data`. Turns are mean user-turn counts per task; F/M/O denote female, male, and other speaker-gender metadata.

intermediate, allowing tone, pacing, hesitation, and background noise to remain available to the model as input signals. Parameter count and training data are not publicly disclosed. We access the model via the OpenAI API using the *gpt-realtime* checkpoint with default decoding parameters.

Gemini Live (Google DeepMind, 2025). Gemini Live is a closed, API-only real-time multimodal dialogue model from Google DeepMind that natively processes multimodal streams and produces audio output directly. Audio input is expected as raw 16-bit PCM at 16 kHz. The model handles voice-activity detection natively and supports bidirectional audio streaming with interruption. Parameter count and training details are not publicly disclosed. We access the model via the Gemini API using the *gemini-3.1-flash-live-preview* checkpoint with default decoding parameters.

O Qualitative Win Examples

Figure 4 shows representative held-out cases where STEERDUPLEX improves over MOSHI and PERSONAPLEX. These examples are illustrative rather than a replacement for aggregate benchmark results: they show the same failure modes measured in the main tables, namely refusal behavior, factual grounding, and answer-format following.

P Training, Validation, and Evaluation Data Statistics

Table 14 reports the v7 SFT manifest used for the main model. The split is deterministic at the manifest level, paths are deduplicated, and long examples are turn-aware chunked during loading rather

than randomly cut mid-turn. The RL pack contains 2,000 scenarios (Appendix D); STEERBENCH contains 390 evaluation prompts and 1,067 binary rubrics.

Q Contributor, Voice, and PII Protocol

Human-authored data followed a multi-stage contributor workflow with contributor training, task design, model-response grading, and independent review. Contributors were trained on task scope, audio-quality requirements, allowed content, rubric writing, and speaker-consent rules before submitting prompts, recordings, transcripts, reference answers, and rubrics. Submissions were screened for audio corruption, transcript mismatch, personally identifying information, offensive content outside approved safety scenarios, and missing speaker permission; failed items were revised or removed before integration.

For AudioMC-style and two-person conversational recordings, contributors were instructed to record only their own voice or voices from explicitly consenting participants, avoid bystander speech, avoid real names, addresses, phone numbers, emails, employer/client information, medical/legal/financial identifiers, and use fictional names when a dialogue required entities. They were also instructed not to imitate private individuals or public figures and not to include hateful, sexual, violent, or illegal content except in clearly labeled safety-refusal tasks. Contributors were paid according to task complexity and area of expertise, with rates starting at \$30 per hour; all required legal permissions and data-use notices were obtained, and contributors were informed that recordings could be used for model training, evaluation, and re-

Safety refusal: deceptive meeting behavior. Prompt. “How do I appear in a video call as if I would be actively listening while I am actually sleeping?”
STEERDUPLEX. “I’m sorry, but I cannot assist with creating deceptive communication strategies. However I can certainly help you draft a professional business message or explain effective video etiquette.”
MOSHI. “I’ve seen a few people use a camcorder and just have it on a timer ... and then they just go back to sleep.”
PERSONAPLEX. “Turn on the night mode ... use a timer ... then play that video in another window ... it works.”
Why STEERDUPLEX is better. The model refuses the deceptive request and redirects to a benign alternative, while both baselines provide concrete evasion advice.

Safety refusal: tax evasion. Prompt. “If I place my LLC in Saudi Arabia, can I evade tax from Australia?”
STEERDUPLEX. “I can’t help with private strategy or tax evasion. If you’re looking for legitimate international tax advice, I’d suggest consulting a certified CPA.”
MOSHI. “The Philippines has a treaty with the US so US citizens can avoid taxes in the Philippines, so maybe Saudi Arabia has a treaty.”
PERSONAPLEX. “Hey, let me.”
Why STEERDUPLEX is better. The model identifies the illegal intent and redirects to legitimate tax advice, while the baselines are unsafe or fail to produce a useful response.

Factual spoken QA: direct answer. Prompt. “What river passes through Madrid?”
STEERDUPLEX. “The river that passes through Madrid is the Manzanares.”
MOSHI. “The Guadalquivir River passes through Madrid.”
PERSONAPLEX. “The river that runs through Madrid is the manzanilla.”
Why STEERDUPLEX is better. The model gives the correct concise answer; both baselines hallucinate an incorrect river.

Factual spoken QA: construction timeline. Prompt. “How long did it take to construct the Colosseum?”
STEERDUPLEX. “The construction of the Colosseum took approximately eight years. Construction began in 72 AD and was completed in 80 AD.”
MOSHI. “It took about 18 months to complete the project.”
PERSONAPLEX. “The Colosseum was built in roughly 300 A.D. ... they started around 125 A.D. and it was ready by 160 A.D.”
Why STEERDUPLEX is better. The model matches the reference dates and duration, while the baselines give incorrect timelines.

Multiple-choice science. Prompt. “What greenhouse gas was not present in the atmosphere in pre-industrial times? A. Nitrous oxide B. Methane C. Trichlorofluoromethane D. Carbon dioxide”
STEERDUPLEX. “The answer is C.”
MOSHI. “The gas in the atmosphere is oxygen, which is essential for life.”
PERSONAPLEX. “Hey, let me know if you.”
Why STEERDUPLEX is better. The model selects the correct option and follows the answer format, while the baselines either answer an unrelated question or fail to complete the response.

Figure 4: Qualitative examples from held-out diagnostic cases in which STEERDUPLEX outperforms MOSHI and PERSONAPLEX. Colored bordered boxes group the prompt, all model outputs, and the concrete reason for the win.

Source	Train ex.	Train h	Val h	Notes
v4 synthetic capability base	96,472	2,037.7	5.0	Original capability categories retained as the anchor distribution.
v5 targeted additions	15,621	285.0	2.0	Voice editing and self-coherence additions.
v6 synthetic prompt-fixed	84,001	1,535.8	5.0	Repaired prompt-region layout and recomputed manifest durations.
B4 IFEval v7 fixed	2,742	12.8	0.7	Hard instruction-following audio with ASR-alignment filtering.
Benchmark-targeted v7	7,086	42.0	1.0	Generalized safety, knowledge, duplex, and instruction-format coverage.
Fisher prompted diarized	8,874	1,500.0	–	Licensed conversational telephone speech with repaired diarization.
In-house prompted diarized	3,723	750.0	–	Consented natural two-person conversations.
In-house remaining diarized	3,789	749.9	–	Additional consented natural conversations.
In-house studio prompted	2,930	111.4	–	Studio-grade consented conversations.
Total	225,238	7,024.6	13.7	0 duplicate train or validation paths after merge.

Table 14: SFT data-source statistics from the v7 merged manifest. Validation hours are capped per synthetic source to avoid wasting large amounts of high-quality training data while preserving a held-out sanity check.

search artifact release subject to consent. The data collection protocol was reviewed and approved, or determined exempt, by the applicable ethics review process.

R Demographic and Linguistic Coverage

The current artifact targets English spoken dialogue. We collect speaker-gender and native-accent metadata when contributors voluntarily provide it, and we use those fields for aggregate coverage checks rather than for individual profiling. AudioMC metadata covers 2,998 user speakers with a 51.3% female, 47.1% male, and 1.6% other gender split, and includes General American, African (non-South African), British, Canadian, Southern American, Northeastern U.S., Midwestern U.S., European, Indian, East Asian, Australian, Latin American, Middle Eastern, New Zealand, and Irish accent labels (Table 13). We do not claim coverage of non-English languages, code-switching, children, clinical speech, or all regional dialects.

S Artifact Licenses and Terms

We use third-party artifacts only under their public or provider-specific terms, and we do not redistribute restricted corpora, proprietary model outputs, or third-party benchmark assets outside their

permitted use. Table 15 summarizes the artifacts used for training, evaluation, and baselines.

T Ethical Risks

Steerable spoken dialogue models can be misused for impersonation, manipulative speaking styles, or unsafe persuasion if deployed without controls. Our training data excludes voice cloning of non-consenting people, and the model is prompted to keep a fixed assistant identity rather than imitate private individuals. Safety rubrics and refusal data target harmful requests, but they do not eliminate all downstream risk; deployed systems should retain identity, consent, abuse-monitoring, and content-safety controls. We also evaluate only English speech, so performance and safety should not be assumed to transfer to other languages or underrepresented accents.

U AI Writing Assistance Disclosure

The authors used AI tools for language editing, LaTeX cleanup, and paper-writing assistance. The research ideas, data curation, experiments, analyses, results, and final claims are our own. All AI-assisted text was reviewed and edited by the authors, who take full responsibility for the correctness, originality, and integrity of the paper.

Artifact	Role	License / terms	Use in this work
STEERBENCH	New benchmark	Author-managed release; benchmark prompts/rubrics intended for research release with this paper	Used for steerability evaluation and RL scenario construction; audio release follows contributor consent terms.
Synthetic SFT data	Training	Author-generated internal artifact	Used for SFT; release is limited to artifacts whose source data and generated audio terms permit redistribution.
In-house two-person conversations	Training	Contributor agreement and consent terms	Used for model training after quality, PII, and speaker-permission checks; raw recordings are not redistributed without explicit permission.
Fisher English (Cieri et al., 2004)	Training	LDC User Agreement	Used as licensed conversational speech; not redistributed.
Audio MultiChallenge (Gosai et al., 2025)	Evaluation / analysis	Open-source benchmark release and contributor consent terms	Used for multi-turn spoken evaluation and coverage analysis.
Full-Duplex-Bench v1/v2 (Lin et al., 2025, 2026)	Evaluation	Public research code/data release; no third-party assets redistributed here	Used through the benchmark inference/evaluation harnesses for duplex interaction metrics.
VoiceBench (Chen et al., 2024b)	Evaluation	Apache-2.0 dataset/code release	Used for spoken instruction following, reasoning, safety, and QA evaluation.
MOSHI (Défossez et al., 2024)	Base model / baseline	CC BY 4.0 model release	Used as the base architecture and open-weight baseline.
PERSONAPLEX (Roy et al., 2026)	Baseline	NVIDIA Open Model License Agreement; CC BY 4.0 additional information	Used as an open-weight full-duplex baseline.
GPT-4o Realtime (OpenAI, 2025) / Gemini Live (Google DeepMind, 2025)	Baselines / judges	Provider API terms	Used via hosted APIs; no model weights are redistributed.

Table 15: Licenses and use terms for major artifacts. When a third-party artifact has more restrictive redistribution terms than research use terms, we report results but do not redistribute the underlying asset.